A Head Camera Is (Almost) All You Need

An electronics-free UMI glove labelled and observed by one moving camera

Yosub Shin[1] Yuan Liu[2] Igor Molybog[1] [1]University of Hawai'i at Mānoa [2]Independent

arXiv:2610.27622
The UMI Arena, CoRL 2026

Why a head camera only

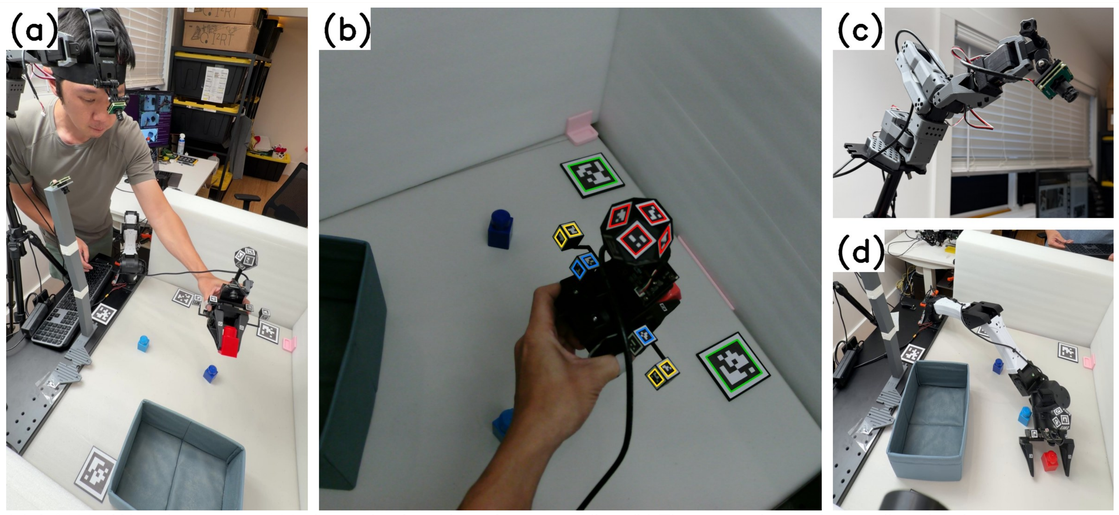
- Wrist devices need a **second stream** to track the gripper: SLAM, an IMU, a controller or an external camera.
- **Gaze leads the hand**: the gripper stays in view, at arm's length, with the goal.
- The frames the policy sees are the frames that **carry the labels**.

No SLAM, no IMU, no wrist camera, no cables.

$144 glove, camera, strap

103 g on the head

One head camera: the policy's observation and the pose labels



(a) A $122 4K head camera, a $22 electronics-free glove, four paper tiles. (b) The camera's own frame: tiles give the camera pose, the ball the gripper pose, the jaw tags the opening. (c) At deployment the same camera rides a $122 six-joint arm (d) beside the robot.

Label accuracy

- **Camera pose** from the paper tiles, refit from the footage to 0.5 px.
- **Gripper pose** by PnP on a 15 mm ArUco dodecahedron.
- **Jaw width** from printed tails on the jaws.

6.5 mm tool point

1.6° orientation

3.5 % jaw width, of stroke

Median error against the robot's encoders; a label on 93–100 % of frames.

Collection

2–3× demos per minute vs teleop

198 demos in four sessions

6–8 s per demo; teleop 14 s

0 robot time

Continuous recording, episodes cut by a keypress; 0.7–7.4 % of frames dropped for tracking.

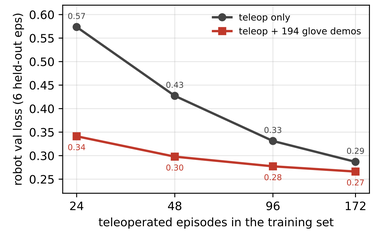
Deployment

- The same camera on a $122 six-joint arm beside the YAM.
- It follows the gripper **from the robot's encoders alone**, in teleoperation and in rollouts, so the policy sees what the collector saw.

46 mm head pose vs the collector

5.4° static camera: 87 mm, 7.0°

Cheap glove demos complement a small teleop set



10/25
48 teleop episodes

17/25
48 teleop + 194 glove demos

1 ≈ 2–3
teleop episode ≈ glove demos

9 placements won to 2 lost ($p = 0.065$). On the curve (head-label chain, own held-out targets) the 194 glove demos are worth **65–100 teleop episodes**. With all 172 they still help: 18/25 → 20/25 (4 won to 2 lost, $p = 0.69$).

Head-camera labels can replace the encoders

teleop labelled by	encoders	head camera
validation loss	0.318	0.328 (+3 %)
rollouts	18/25	15/25 ($p = 0.58$)

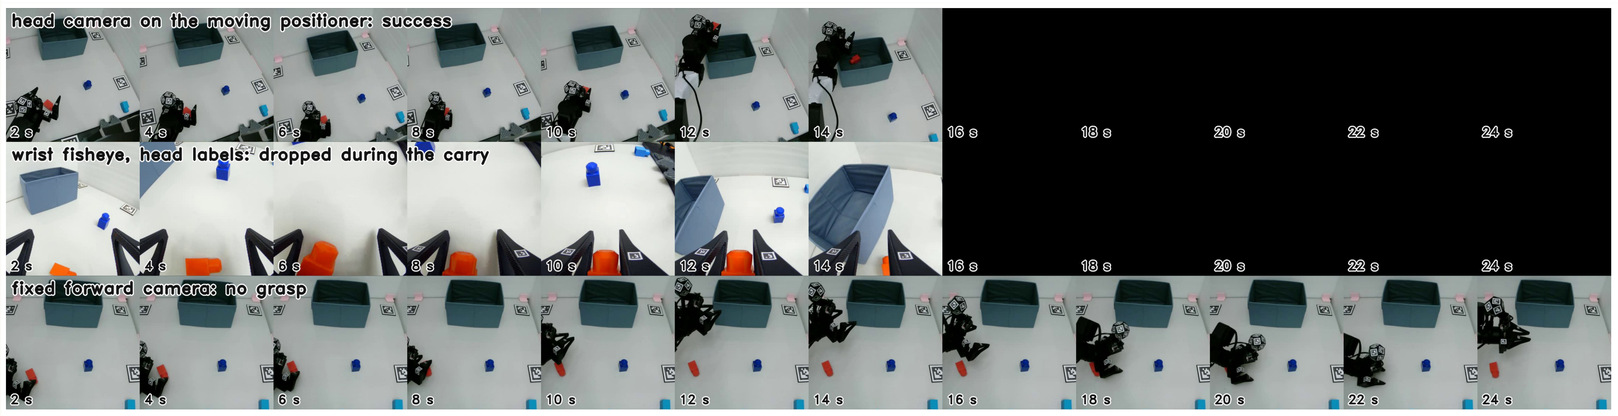
Same 172 teleop episodes, same held-out encoder targets; 5 placements won to 8 lost ($p = 0.58$).

Where each view fails

failures of 25	no grasp	dropped	success
head camera	**5**	0	20
wrist fisheye (UMI)	6	**4**	15
wrist fisheye, head labels	3	**8**	14
fixed camera	**12**	1	12

Wrist policies **drop the block mid-carry**, always while the fisheye looks away from the bin. Head and fixed cameras keep the bin in frame and fail only by not grasping.

One placement, seen by each policy



Placement 23, each policy's own frames every 2 s on one clock, black once its episode ended. Head camera: gripper and bin in one frame through the carry, done in 16 s. Wrist fisheye: only block and wall after the grasp; drops the block before reaching the bin. Fixed camera: sees the whole scene but never closes the gripper (41 s, cut here).

It is the view, not the labels

Same view, other labels
wrist fisheye observation

glove labels from	success	dropped
fixed camera	15/25	4
head camera	14/25	8

Same labels, other view
head-camera glove labels

observation	success	dropped
wrist fisheye	14/25	8
head camera	**20/25**	**0**

Limits and next

Limits
- One task, one collector, one tile layout.
- Teleop watched the head-camera view.
- Glove-only policies do not transfer; co-training is required.

Next
- Precision tasks, e.g. cable ties.
- Bimanual, with a second glove.
- Tile-free sessions in the wild.

Setup: eight ACT policies, joint-space actions. Rollouts: 25 placements, one trial each, 60 s. Every co-trained mix uses the same 194 glove demos.